

AN INTRODUCTION TO HIGH-THROUGHPUT COMPUTING AT ISYE

CREATING THE NEXT®

WHAT IS COVERED IN THIS TUTORIAL



- How to get access to the ISYE Cluster
- How to get your files to the ISYE Cluster
- How to access the Cluster
- What is HTCondor ?
- What the different between HPC and High-Throughput Computing
- How to Run a Job with HTCondor
- How to Monitor a HTCondor job
- Submitting a simple jobs with HTCondor
- Submitting Multiple Jobs with HTCondor
- Testing and Troubleshooting
- Cluster Documentation

- ISyE
 - HTCondor Cluster – Group of Linux servers with standard set of software and networked home drives scheduled via HTCondor and maintained by ISyEIT staff and housed within Groseclose. **Open to all members of the ISyE community** (faculty, undergraduate, graduate, & postdocs)

GETTING ACCESS TO THE ISYE CLUSTER

- First, sign up for an ISyE account: <https://www.isye.gatech.edu/about/school/computing>
- Get a terminal app:
 - Windows:
 - PuTTY: <https://www.chiark.greenend.org.uk/~sgtatham/putty/> (careful if you Google!)
 - SecureCRT: <https://software.oit.gatech.edu>
 - Windows Subsystem for Linux: <https://docs.microsoft.com/en-us/windows/wsl/install-win10>
 - Mac:
 - Terminal.app (built-in)
 - iTerm2: <https://www.iterm2.com>
 - Linux: It's built in!
 - iOS: Prompt: <https://panic.com/prompt>
 - Android: lots of options <https://www.dunebook.com/10-best-android-terminal-emulator/>

GETTING YOUR FILES TO THE CLUSTER



File access via SCP/SFTP:

- Windows:
 - WinSCP: <https://winscp.net>
 - FileZilla: <https://filezilla-project.org/>
- Mac:
 - Cyberduck: <https://cyberduck.io>
 - FileZilla: <https://filezilla-project.org>
 - Transmit: <https://panic.com/transmit> (not free)
- Linux:
 - scp at the command line, GUI varies by distro
- GitHub

HOW TO ACCESS THE ISYE CLUSTER



Off campus, or on GTwireless:

- Connect to VPN (<https://faq.oit.gatech.edu/content/how-do-i-get-started-campus-vpn>)

OR

- SSH first to **castle.isye.gatech.edu**

Once connected to VPN or castle, or if you're on the wired ISyE network:

- SSH to **wren.isye.gatech.edu**
- SCP/SFTP via **castle.isye.gatech.edu** (does not require VPN or an SSH connection to castle first)

- Documentation:
 - ISyE website (better & more documentation coming soon!)
 - <https://www.isye.gatech.edu/about/school/computing>
 - <https://www2.isye.gatech.edu/apps/helpdesk/kb/browse/866-Computational-Resources>
 - HTCondor website (make sure you select the 'stable' release)
 - <https://research.cs.wisc.edu/htcondor/manual/>
 - <https://research.cs.wisc.edu/htcondor/manual/quickstart.html>
 - <http://chtc.cs.wisc.edu/multiple-jobs.shtml>

Open source software for distributed High-Throughput computing. Developed at the University of Wisconsin in the 1980's. HTCondor has three components

- A job scheduler
- A resource manager
- A workflow management system

High-Throughput Computing - Allows for many computational tasks to be completed over a long period of time. Useful for researchers and other users who are more concerned with the number of computations they can do over long spans of time than they are with short-burst computations


HTCondor manages and runs work on your behalf
Schedule tasks on a single computer to not overwhelm the computer

Schedule tasks on a group* of computers (which may/may not be directly accessible to the user)

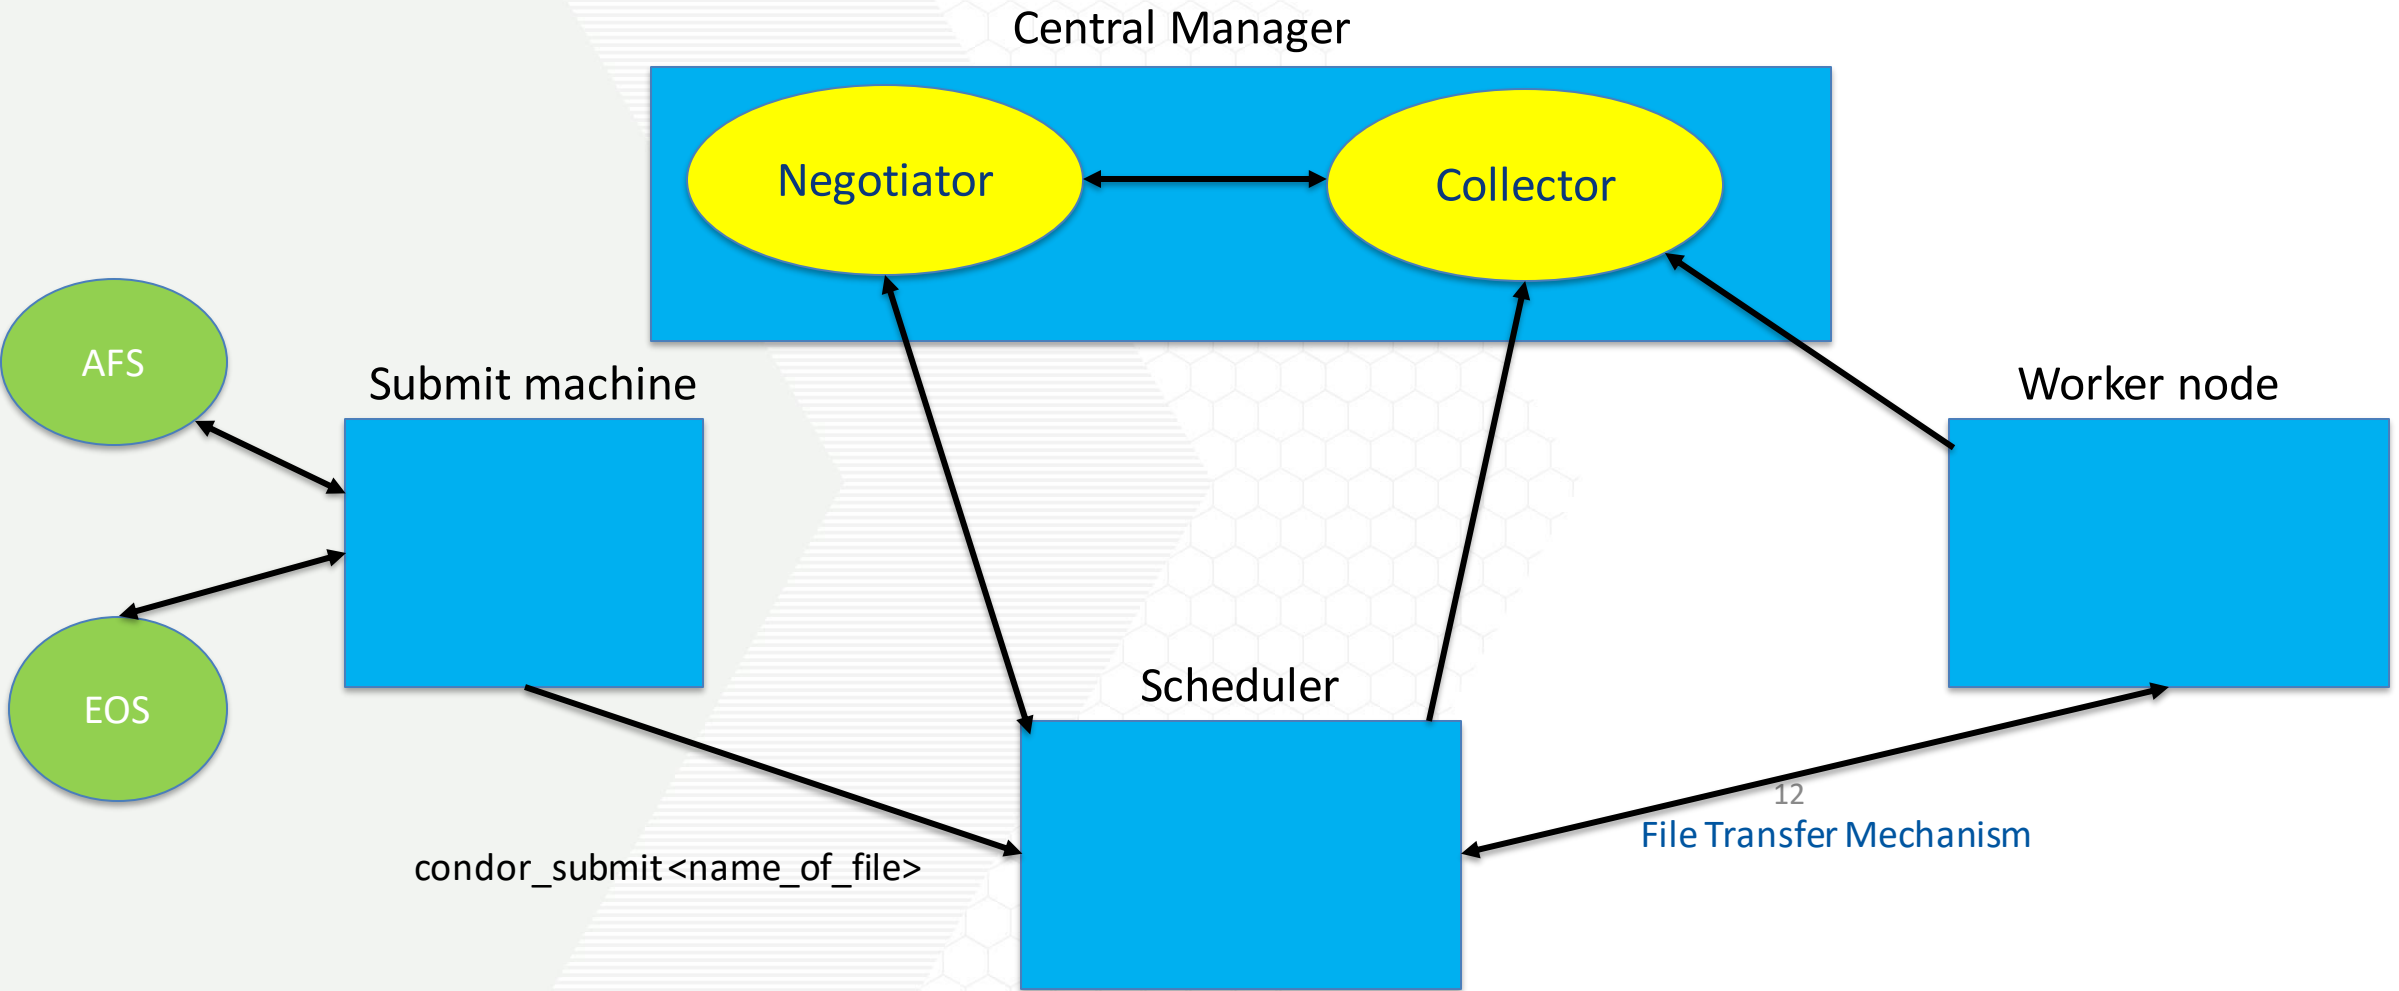
Schedule tasks submitted by multiple users on one or more computers

- HTCondor has built-in ways to submit multiple independent jobs with one submit file.
 - Analyze multiple data files
 - Test parameter or various input combination
- ...without having to:
 - Start each job individually
 - create separate submit files for each job

HTCONDOR WORKFLOW



HTCONDOR SERVICE COMPONENTS



- A single computing task is called a “job”
- Three main pieces of a job are the input, executable (program) and output
- Executable must be runnable from the command line without any interactive input

- Jobs are nearly always using a part of a computer, not the whole thing
- Very important to request appropriate resources
 - memory,
 - cpus
 - gpus
 - disk
- Even if your system has default CPU, memory and disk requests, these may be too small!
- Important to run test jobs and use the log file to request the right amount of resources
- Requesting too little: causes problems for your and other jobs; jobs might be held by HTCondor and requesting too much: jobs will match to fewer “slots”

Simple Submit File



```
universe = vanilla  
  
getenv = true  
  
executable = lp.py  
  
arguments = test0.lp  
  
Log = test.log  
  
output = test.out  
  
error = test.error  
  
notification = error  
  
notification = complete  
  
request_memory = 1024  
  
request_cpus = 1  
  
queue
```

HOW TO SUBMITTING AND MONITOR YOUR JOB

To submit a job/jobs:

condor_submit *submit_file_name*

To monitor submitted jobs, use:

condor_q

condor_q -nobatch

- Show the user's job only
- Jobs summarized in "batches"
- -nobatch to see extra detail about what is happening with a job.

- Jobs can go wrong “internally”:
 - something happens after the executable begins to run
- Has a badly formatted executable
- Uses too much memory
- Has a badly formatted command file
- Asked for too many resources.

TROUBLESHOOTING COMMANDS



`condor_q -better <jobid>`

`condor_q -better <jobid> -machine <name>`

`condor_q -better -reverse -machine <machine>`

`condor_q -hold`

- To submit a job/jobs:

condor_submit *submit_file_name*

```
$ condor_submit mul.cmd
Submitting job(s).
1 job(s) submitted to cluster 102329.
```

- To monitor submitted jobs, use:

condor_q

```
$ condor_q
```

```
-- Schedd: isye-hps0005.isye.gatech.edu: <10.1.112.2:9618?... @ 03/04/20 13:11:06
OWNER  BATCH_NAME  SUBMITTED  DONE  RUN  IDLE  TOTALJOB_IDS
kross48 ID: 102329  3/4  13:10   _   1   _   1 102329.0
```

Total for query: 1 jobs; 0 completed, 0 removed, 0 idle, 1 running, 0 held, 0 suspended

Total for kross48: 1 jobs; 0 completed, 0 removed, 0 idle, 1 running, 0 held, 0 suspended

Total for all users: 22 jobs; 0 completed, 0 removed, 1 idle, 21 running, 0 held, 0 suspended

Idle (I)	<ul style="list-style-type: none">• The job is waiting in the queue to be executed
Running (R)	<ul style="list-style-type: none">• The job is running
Completed (C)	<ul style="list-style-type: none">• The job is exited
Held (H)	<ul style="list-style-type: none">• Something is wrong with the submission file• The user executed <code>condor_hold <job_id></code>
Suspended (S)	<ul style="list-style-type: none">• The user executed <code>condor_suspend</code>
Removed (X)	<ul style="list-style-type: none">• The job has been removed by the user.

OTHER USEFUL COMMANDS



- `condor_status -master` – View all compute nodes and their resources
- `condor_q -ana:sum (user)`
- `condor_rm <jobid>` - Remove a job
- `condor_qedit <jobid> RequestMemory 3078` – Edit an existing job
- `condor_q -all` – See the whole queue (All users, All Jobs)
- `condor_rm -a` - Removes all of your jobs.
- `condor_q -af HoldReason`
- `condor_q -analyze 107.5`
- `condor_status -compact -constraint 'TotalGpus > 0'` - View available GPU systems
- `condor_status -compact -constraint 'TotalGpus > 0' -af Machine TotalGpus CUDADeviceName CUDACapability` – Display GPU hardware information

-
- Jobs are nearly always using a part of a computer, not the whole thing
 - Very important to request appropriate resources (memory, cpus, disk) for a job

Submitting Multiple jobs Using \$(ProcId)

```
universe = vanilla
```

```
getenv = true
```

```
executable = $ENV(HOME)/gurobi/lp.py
```

```
arguments = test$(ProcId).lp
```

```
Log = $ENV(HOME)/gurobi/$(Cluster).log
```

```
output = $ENV(HOME)/gurobi/$(Cluster).$(ProcId).out
```

```
error = $ENV(HOME)/gurobi/$(cluster).$(ProcId).error
```

```
notification = error
```

```
notification = complete
```

```
request_memory = 1048
```

```
request_cpus = 1
```

```
queue 3
```

Use the `$(ClusterId)`, `$(ProcId)` variables to provide unique values to jobs.*

Submitting Multiple jobs Using Queue Statements

```
universe = vanilla
```

```
getenv = true
```

```
executable = lp.py
```

```
arguments = $(file)
```

```
Log = $ENV(HOME)/gurobi/$(Cluster).log
```

```
output = $ENV(HOME)/gurobi/$(Cluster).$(process).out
```

```
error = $ENV(HOME)/gurobi/$(cluster).$(Process).error
```

```
notification = error
```

```
notification = complete
```

```
request_memory = 1048
```

```
request_cpus = 1
```

```
queue file from testlp.txt
```

Submitting Multiple jobs Using InitialDir

```
universe = vanilla  
  
getenv = true  
  
initialdir = job$(process)  
  
executable = lp.py  
  
arguments = test.lp  
  
Log = job.log  
  
output = job.out  
  
error = job.error  
  
notification = error  
  
notification = complete  
  
request_memory = 1048  
  
queue 3
```

APPLICATIONS AVAILABLE ON THE CLUSTER



- Ampl
- Mathematica
- Matlab
- Openmpi
- Cplex
- Python3
- Gams
- RH
- Gurobi
- xpressmp
- Julia
- Lua

`periodic_remove = JobStatus == 2 && (CurrentTime - EnteredCurrentStatus > 24*60*60)`

`periodic_remove = JobStatus == 2 && (CurrentTime - EnteredCurrentStatus > 24*60*60)`

HTCONDOR OUTPUT AND LOGS

These files are defined in the submit file

Output

The STDOUT of the command or script

Error

The STDERR of the command or script

Log

Information about job's execution

execution host

the number of times this job has been submitted

the exit value, etc

Can use relative or absolute paths for all of them

HTCondor will search for this directory

So it should be already created

VIEW MACHINES ATTRIBUTES



```
[kross48@isye-hps0005 examples]$ condor_status--compact
```

Machine	Platform	Slots	Cpus	Gpus	TotalGb	FreCpu	FreeGb	CpuLoad	ST	Jobs/Min	MaxSlotGb
isye-abel.isye.gatech.edu	x64/RedHat7	0	32		377.71	32	377.71	0.00	Ui	0.00 *	
isye-ames.isye.gatech.edu	x64/RedHat7	0	24		125.68	24	125.68	0.00	Ui	0.00 *	
isye-boyle1.isye.gatech.edu	x64/RedHat7	1	16		47.00	15	15.75	0.06	**	0.00	31.25
isye-boyle10.isye.gatech.edu	x64/RedHat7	0	16		47.00	16	47.00	0.00	Ui	0.00 *	
isye-bulldozer.isye.gatech.edu	x64/RedHat7	1	32		125.73	31	121.73	0.03	**	0.00	4.00
isye-dollar.isye.gatech.edu	x64/RedHat7	1	12		39.12	11	7.87	0.08	**	0.00	31.25
isye-fisher1.isye.gatech.edu	x64/RedHat7	0	24		125.73	24	125.73	0.00	Ui	0.00 *	
isye-fisher2.isye.gatech.edu	x64/RedHat7	0	24		125.73	24	125.73	0.00	Ui	0.00 *	
isye-fisher3.isye.gatech.edu	x64/RedHat7	0	24		125.73	24	125.73	0.00	Ui	0.00 *	
isye-fisher4.isye.gatech.edu	x64/RedHat7	0	24		125.73	24	125.73	0.00	Ui	0.00 *	
isye-goddard1.isye.gatech.edu	x64/RedHat7	1	12		70.62	11	6.62	0.13	**	0.00	64.00
isye-goddard2.isye.gatech.edu	x64/RedHat7	1	12		70.62	11	66.62	0.08	**	0.00	4.00
isye-gpu1001.isye.gatech.edu	x64/RedHat7	1	72	8	503.10	71	483.10	0.00	*i	2.00	20.00
isye-hpc0009.isye.gatech.edu	x64/RedHat7	1	10		57.44	9	55.44	0.85	**	0.00	2.00

- Can also combine with
success_exit_code = < Integer >
retry_until = < Integer | Expression >

```
executable = foo.exe  
max_retries = 5  
retry_until = ExitCode >= 0  
queue
```

- If you know your job has a problem and it hasn't yet completed, you can:
 - Place it on hold yourself, with `condor_hold [U/C/J]`

```
$ condor_hold bob  
All jobs of user "bob" have been held
```

```
$ condor_hold 128  
All jobs in cluster 128 have been held
```

```
$ condor_hold 128.0  
Job 128.0 held
```

- Remove it from the queue, using `condor_rm [U/C/J]`

- Jobs that use multiple cores on a single computer can be run in the vanilla universe (parallel universe not needed):


```
request_cpus = 16
```

- If there are computers with GPUs, request them with:

```
request_gpus = 1
```

- HTCondor has different “universes” for running specialized job types
- Vanilla (default)
 - good for most software
- Set in the submit file using:

```
universe = vanilla
```



- A job's log, output and error files can provide valuable information for troubleshooting

Log	Output	Error
<ul style="list-style-type: none">• When jobs were submitted, started, and stopped• Resources used• Exit status• Where job ran• Interruption reasons	Any “print” or “display” information from your program	Captured by the operating system


- Problem: a small number of jobs fail with a known error code; if they run again, they complete successfully.
- Solution: If the job exits with the error code, leave it in the queue to run again

```
max_retries = 3
```

```
job.submit  
  
executable = analyze.exe  
arguments = file$(ProcId).in file$(ProcId).out  
should_transfer_files = YES  
transfer_input_files = file$(ProcId).in  
when_to_transfer_output = ON_EXIT  
  
log = job_$(ClusterId).log  
output = job_$(ClusterId)_$(ProcId).out  
error = job_$(ClusterId)_$(ProcId).err  
  
queue 3
```

- Use the `$(ClusterId)`, `$(ProcId)` variables to provide unique values to jobs.*

* May also see `$(Cluster)`, `$(Process)` in documentation



Automatic Variables

- Each job's ClusterId and ProcId numbers are saved as job attributes
- They can be accessed inside the submit file using:
 - \$(ClusterId)
 - \$(ProcId)

- Create sub-directories* and use paths in the submit file to separate input, error, log, and output files.



* must be created before the job is submitted

- Change the submission directory for each job using `initialdir`
- Allows the user to organize job files into separate directories.
- Use the same name for all input/output files
- Useful for jobs with lots of output files



SEPARATE JOBS WITH initialdir

(submit_dir) /

	job0/	job1/	job2/
job.submit			
analyze.exe	file.in	file.in	file.in
	job.log	job.log	job.log
	job.err	job.err	job.err
	file.out	file.out	file.out

job.submit

```
executable = analyze.exe
initialdir = job$(ProcId)
arguments = file.in file.out
transfer_input_files = file.in
```

```
log = job.log
error = job.err
```

```
queue 3
```

Executable should be in the directory with the submit file, *not* in the individual job directories

(submit_dir)/

job.submit	file0.out	input /	log /	err /
analyze.exe	file1.out	file0.in	job0.log	job0.err
	file2.out	file1.in	job1.log	job1.err
		file2.in	job2.log	job2.err

job.submit

```
executable = lp.py
arguments = file$(Process).in file$(ProcId).out
transfer_input_files = input/file$(ProcId).in

log = log/job$(ProcId).log
error = err/job$(ProcId).err

queue 3
```

- Both the “`from`” and “`in`” syntax support using multiple variables from a list.

`job.submit`

```
executable = compare_states
arguments = -year $(option) -input
$(file)
```

```
should_transfer_files = YES
when_to_transfer_output = ON_EXIT
transfer_input_files = $(file)
```

`job_list.txt`

```
wi.dat, 2010
wi.dat, 2015
ca.dat, 2010
ca.dat, 2015
ia.dat, 2010
ia.dat, 2015
```

```
queue file,option from job_list
```

- Edit executable:
 - Atomically save intermediate states to a checkpoint file
 - Always check for a checkpoint file when starting
- Add HTCondor option that a) saves all intermediate/output files from the interrupted job and b) transfers them to the job when HTCondor runs it again

```
when_to_transfer_output = ON_EXIT_OR_EVICT
```

- By default, a job that is interrupted will start from the beginning if it is restarted.
- It is possible to implement self-checkpointing, which will allow a job to restart from a saved state if interrupted.
- Self-checkpointing is useful for very long jobs, and being able to run on opportunistic resources.

HOW TO INSTALL PYTHON PACKAGES LOCALLY



Procedures to install Python Packages not provided by Default

- `pip install <module> --user`
- `pip install virtualenv --user`

- Remember to use full paths for all files!
 - Including executable directive and all scripts, data files, log files, and output
 - In condor submit file, `$ENV(HOME)=/home/username`
 - In your code `~/` `=/home/username`
 - Not just in condor submission file, but in your code, too!

- In Person: Groseclose 302
- Email: helpdesk@isye.gatech.edu

(please don't email a specific person directly, use the helpdesk email to keep from getting lost!)

ISyE HPC Announcements Email List

<https://lists.isye.gatech.edu/mailman/listinfo/isye-hpc-announce>